



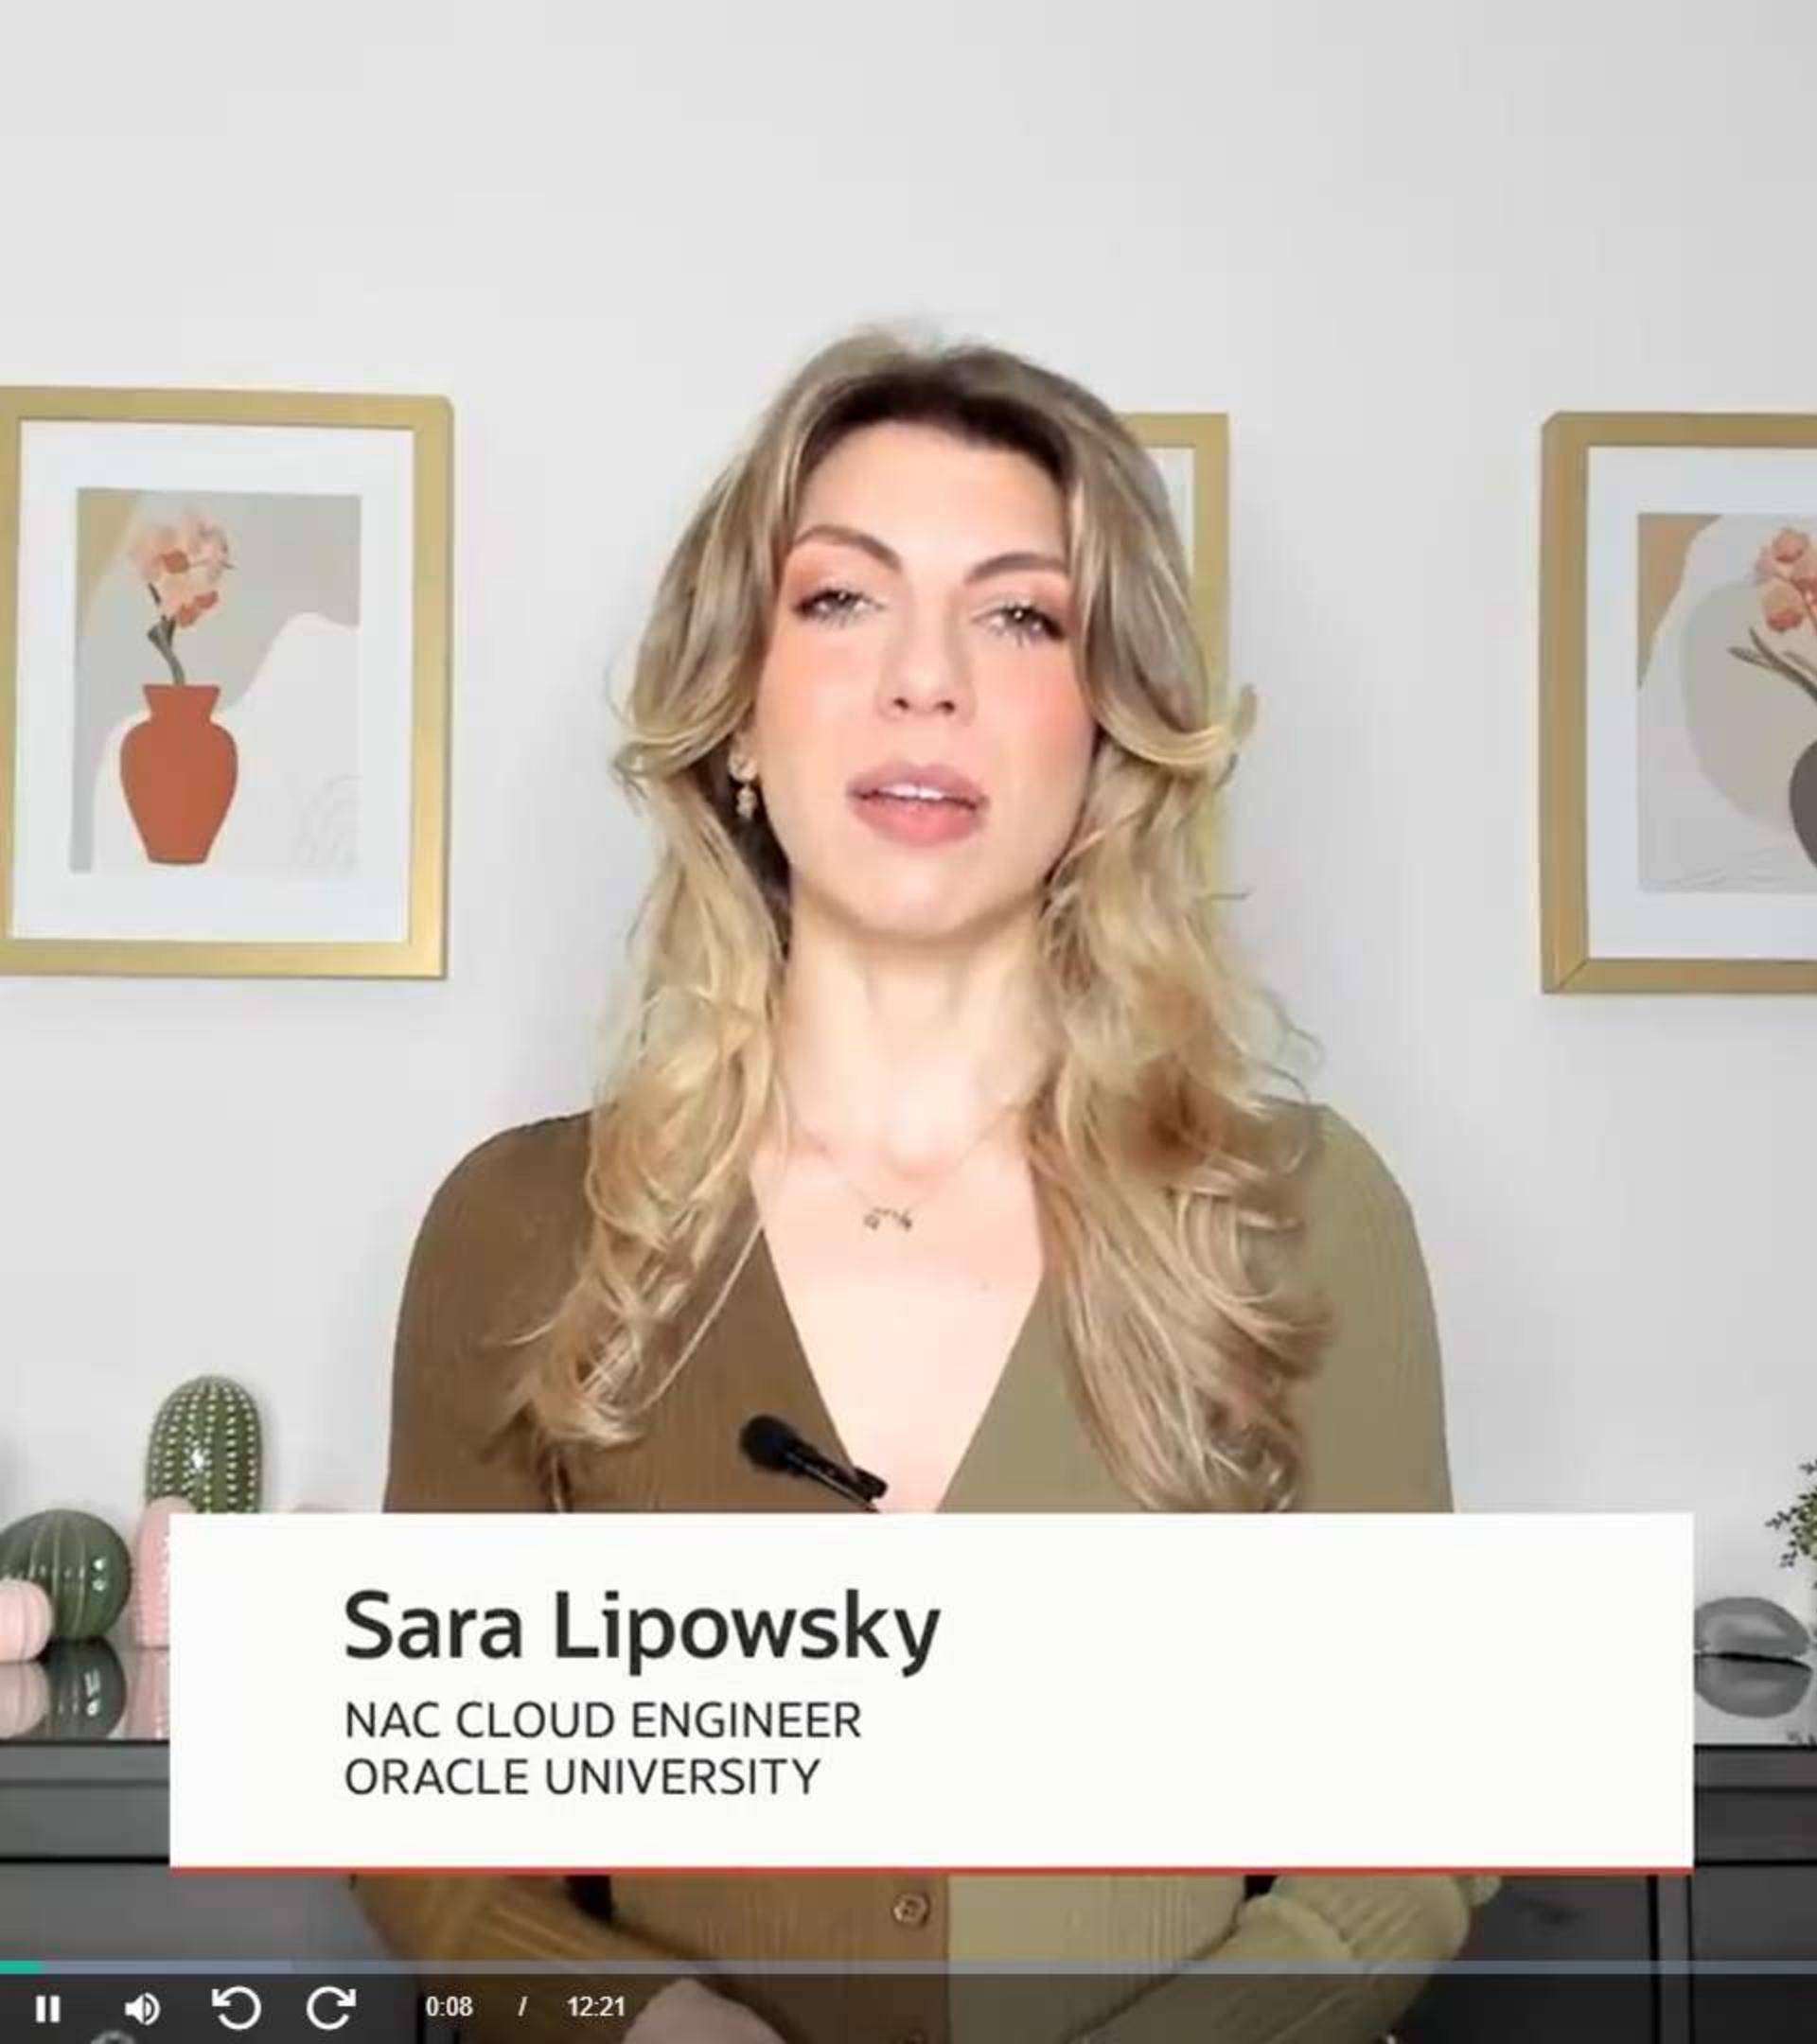
Oracle Cloud Infrastructure

Oracle Autonomous JSON Database

An Overview

Sara Lipowsky

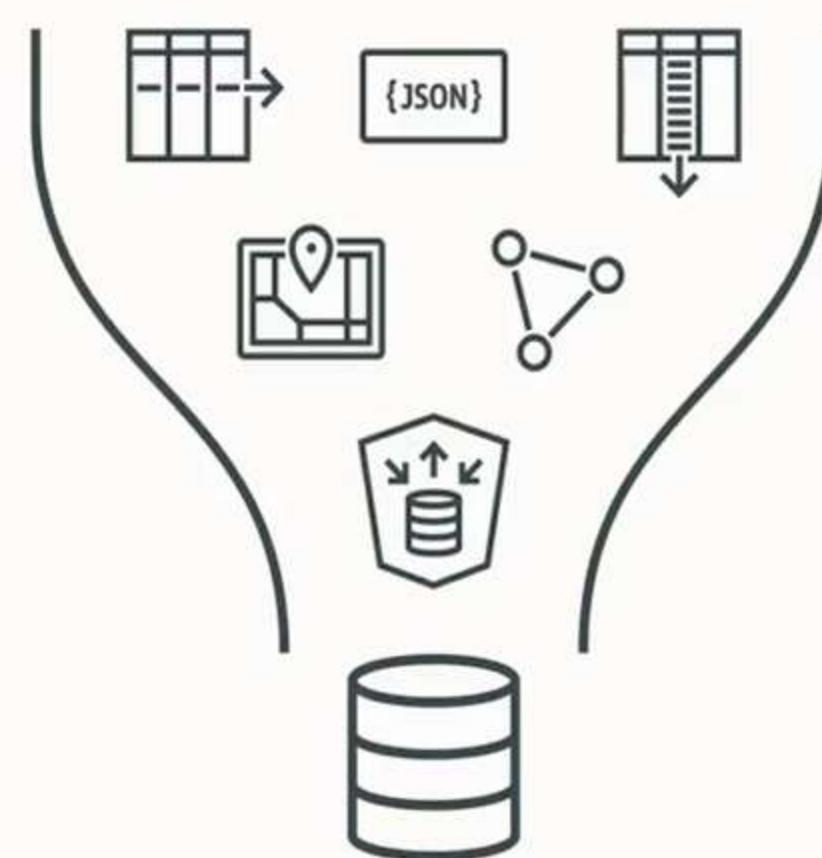
NAC CLOUD ENGINEER
ORACLE UNIVERSITY



Converged Database

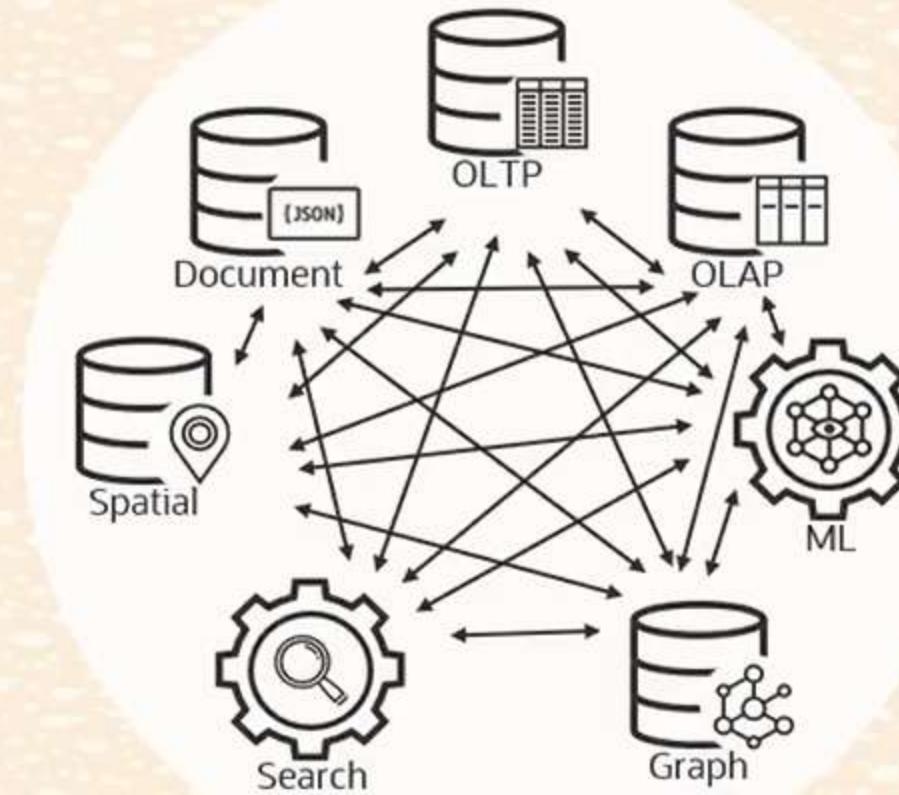


Converged Database Architecture



for **any** data type or workload

Single-purpose databases



for each data type and workload

Multiple security models, query languages, skills, licenses, and so on

Why JSON?

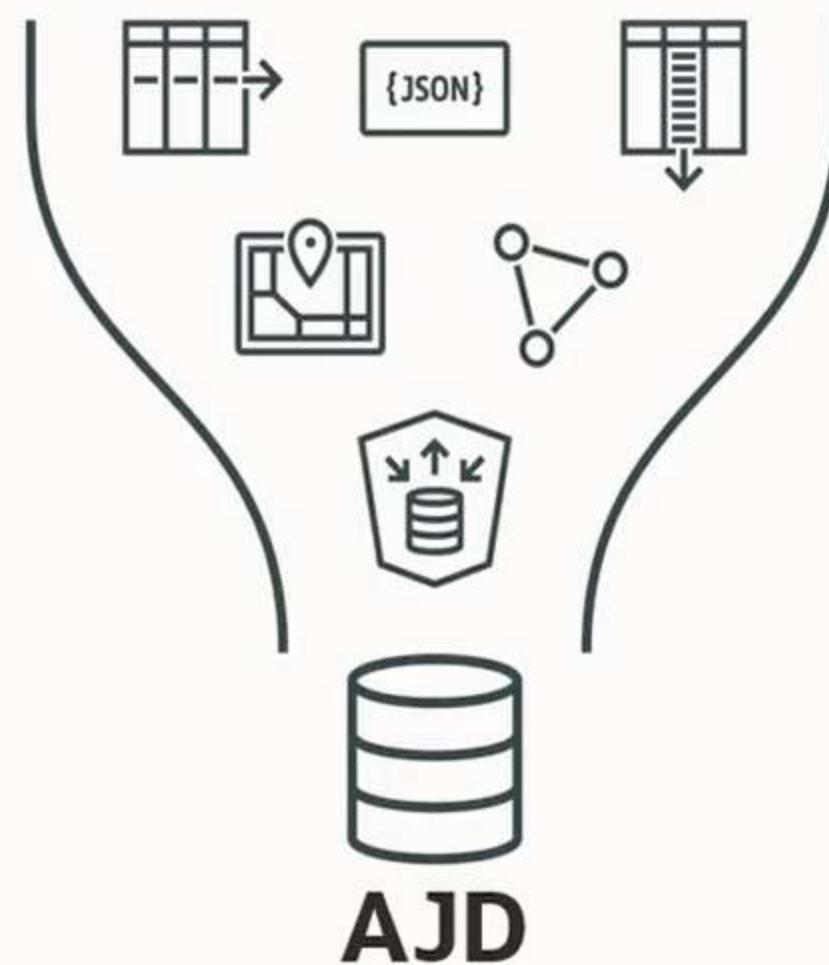
- Schema-flexible
 - Applications can evolve to store new attributes without modifying data definitions.
- Easily consumed by applications
 - Nested structures
 - Maps to application objects
 - Read/write without joins
- Good common format
 - Supported by most programming languages
 - Human readable
 - Simplifies data exchange across app, servers, and database tiers

```
{  
  "name" : "Thomas Anderson",  
  "job"   : "Programmer",  
  "addresses" : [  
    {  
      "street" : "123 Main",  
      "city"   : "Santa Cruz",  
      "zip"    : 95041  
    }  
  ]  
}
```

Autonomous JSON Database

- Low-latency, scalable, JSON storage
- MongoDB APIs or SQL
- No database management
- Always-free service
- All the features of the Autonomous Database
 - Limited to **20 GB for non-JSON** data

The converged database
as a **managed cloud service**



Autonomous JSON Database

Powered by Autonomous Database, a family of Cloud Services introduced in 2018



Self-Driving

Automates all database and infrastructure management, monitoring, tuning



Self-Securing

Protects from both external attacks and malicious internal users



Self-Repairing

Protects from all down time including planned maintenance

Spend Less, Reduce Risk, Innovate More

Autonomous JSON Database

All the benefits of a “one-trick” document database



**Elastic compute
and storage**



**Single-digit latency
reads and writes**



Highly available



**Low price,
always-free tier**

Autonomous Database Workloads



OLTP+ JSON workloads

Transaction Processing (ATP)

Preconfigured for row format, indexes, and data caching to accelerate **transaction processing** and **mixed workloads**

- **Converged database** with no data storage limits (relational or JSON), JSON data fully supported
- Includes **Oracle Database API for MongoDB**



JSON-centric workloads

JSON (AJD)

Price optimized for transactions and analytics on **JSON data**

- Same features as ATP but **75% lower price**
- **Unlimited JSON Collections** + up to 20GB non-JSON data
- Includes **Oracle Database API for MongoDB**



Analytics workloads including JSON data

Data Warehousing (ADW)

Preconfigured for columnar format, partitioning, and large joins to accelerate **analytics**, **data warehouse**, and **data lakehouse**

- Same features as ATP, optimized for analytics
- JSON data fully supported

Single-click upgrade to ATP

Autonomous JSON Database: Pricing and Performance

Autonomous JSON Database pricing:

- \$0.1344 per vCPU-hour, available in 2 vCPU increments (\$240/month for 2 vCPUs)
-

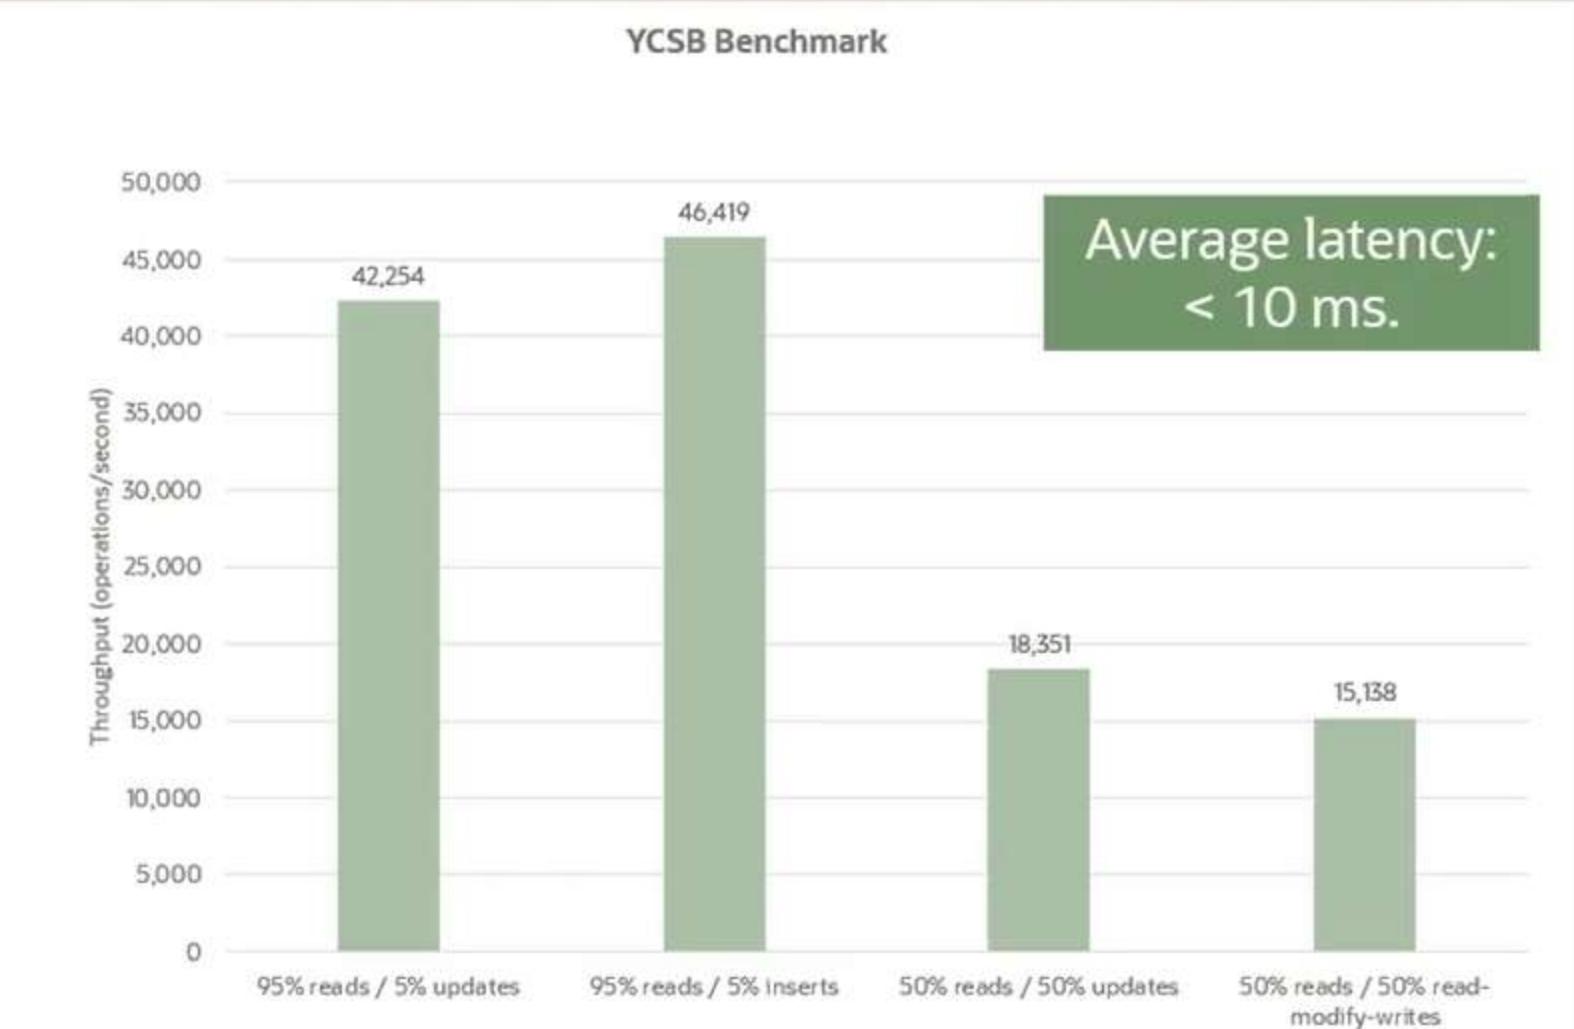
	Autonomous JSON Database	MongoDB Atlas
Configuration	16 vCPUs 1 TB storage	M60 on AWS 16 vCPU 320 GB storage
Price	\$2.74 / hour	\$3.95 / hour*

PLUS: Autonomous JSON Database is auto-scaling; not limited to fixed shapes.

* <https://www.mongodb.com/pricing>

2 vCPUs = 1 OCPU, billing is done on a per-second basis based on OCPUs

Autonomous JSON Database provides consistent performance across different workloads.**



** Based on Autonomous JSON Database with 8 OCPU running in San Jose region

Classic Relational Model

- There are multiple "flat" tables that are related.
- A **table** contains **rows**.
- A **schema** contains tables.
- Rows are structured (the 'S' in SQL).
- Data is accessed using SQL.
- Related entities are joined.

Normalized Tables

id	name	job
123	Anderson	Programmer
345	Smith	Agent

Accessed with SQL

```
select e.id, e.name  
from employee e  
where e.job = 'Programmer'
```

Oracle Database API for MongoDB

- **Data model:** JSON collections, not tables
- **Developers keep their skills** and continue to use MongoDB's tool, drivers, and so on.
- **Easy migrations** of MongoDB workloads to Oracle
- **Enables SQL:**
 - More and faster analytical capabilities, machine learning
 - Query JSON alongside other data models: relational, XML, spatial, and so on
 - Expose relational data, reports, query results as MongoDB collections

Document Collections

- A **document** is a JSON value:
 - Structure is flexible
 - Have a unique **key** (`_id`)
- A **collection** contains documents:
 - Supports insert, get, update, filter
- A **database** contains collections.
- No SQL is required.

MongoDB Collection API

```
use admin;

db.createCollection("employee");

db.employee.insertOne(
{
  "_id" : 123,
  "name" : "Thomas Anderson",
  "job" : "Programmer"
});

db.employee.find(
  {"job" : "Programmer"}
);
```



Oracle API for MongoDB

Database => Schema

Collections created in database “admin” will be in the “ADMIN” schema.

```
use admin;

db.createCollection("employees");

db.employee.insertOne(
{
    "_id" : 123,
    "name" : "Thomas Anderson",
    "job" : "Programmer"
});

db.employee.find(
    {"job" : "Programmer"}
);
```

Oracle API for MongoDB



Collection => Table

Collections are an abstraction or a view of a table with a single JSON column.

```
CREATE TABLE employee (
    ID VARCHAR2,
    DATA JSON
)
```

```
use admin;

db.createCollection("employees");

db.employee.insertOne(
{
    "_id" : 123,
    "name" : "Thomas Anderson",
    "job"   : "Programmer"
});

db.employee.find(
    {"job" : "Programmer"}
);
```

Oracle API for MongoDB



Document => Row

Inserting a document into a collection inserts into the backing table.

```
INSERT INTO employees (data)  
VALUES (:1);
```

```
use admin;  
  
db.createCollection("employee");  
  
db.employee.insertOne(  
{  
    "_id" : 123,  
    "name" : "Thomas Anderson",  
    "job" : "Programmer"  
});  
  
db.employee.find(  
    {"job" : "Programmer"}  
);
```

Oracle API for MongoDB



Filter => Query

Filter expressions are executed as SQL over the backing table; fully utilizes core Oracle Database features such as indexing, cost-based optimization, and so on.

```
SELECT data  
FROM employee e  
WHERE e.data.job = 'Programmer'
```

```
use admin;  
  
db.createCollection("employee");  
  
db.employee.insertOne(  
{  
    "_id" : 123,  
    "name" : "Thomas Anderson",  
    "job" : "Programmer"  
});  
  
db.employee.find(  
    {"job" : "Programmer"}  
);
```

SQL Only When It Is Needed



Oracle API for MongoDB

```
employee.insertOne({  
    "name" : "Bond",  
    "job"  : "Agent"  
});
```

employee

```
{  
    {  
        {  
            "name" : "John"  
            "job"  : "Developer"  
        }  
    }  
}
```

SQL for JSON

```
SELECT  
t.data.name.string(),  
FROM employee t  
WHERE t.data.job.string()  
= 'Agent';
```

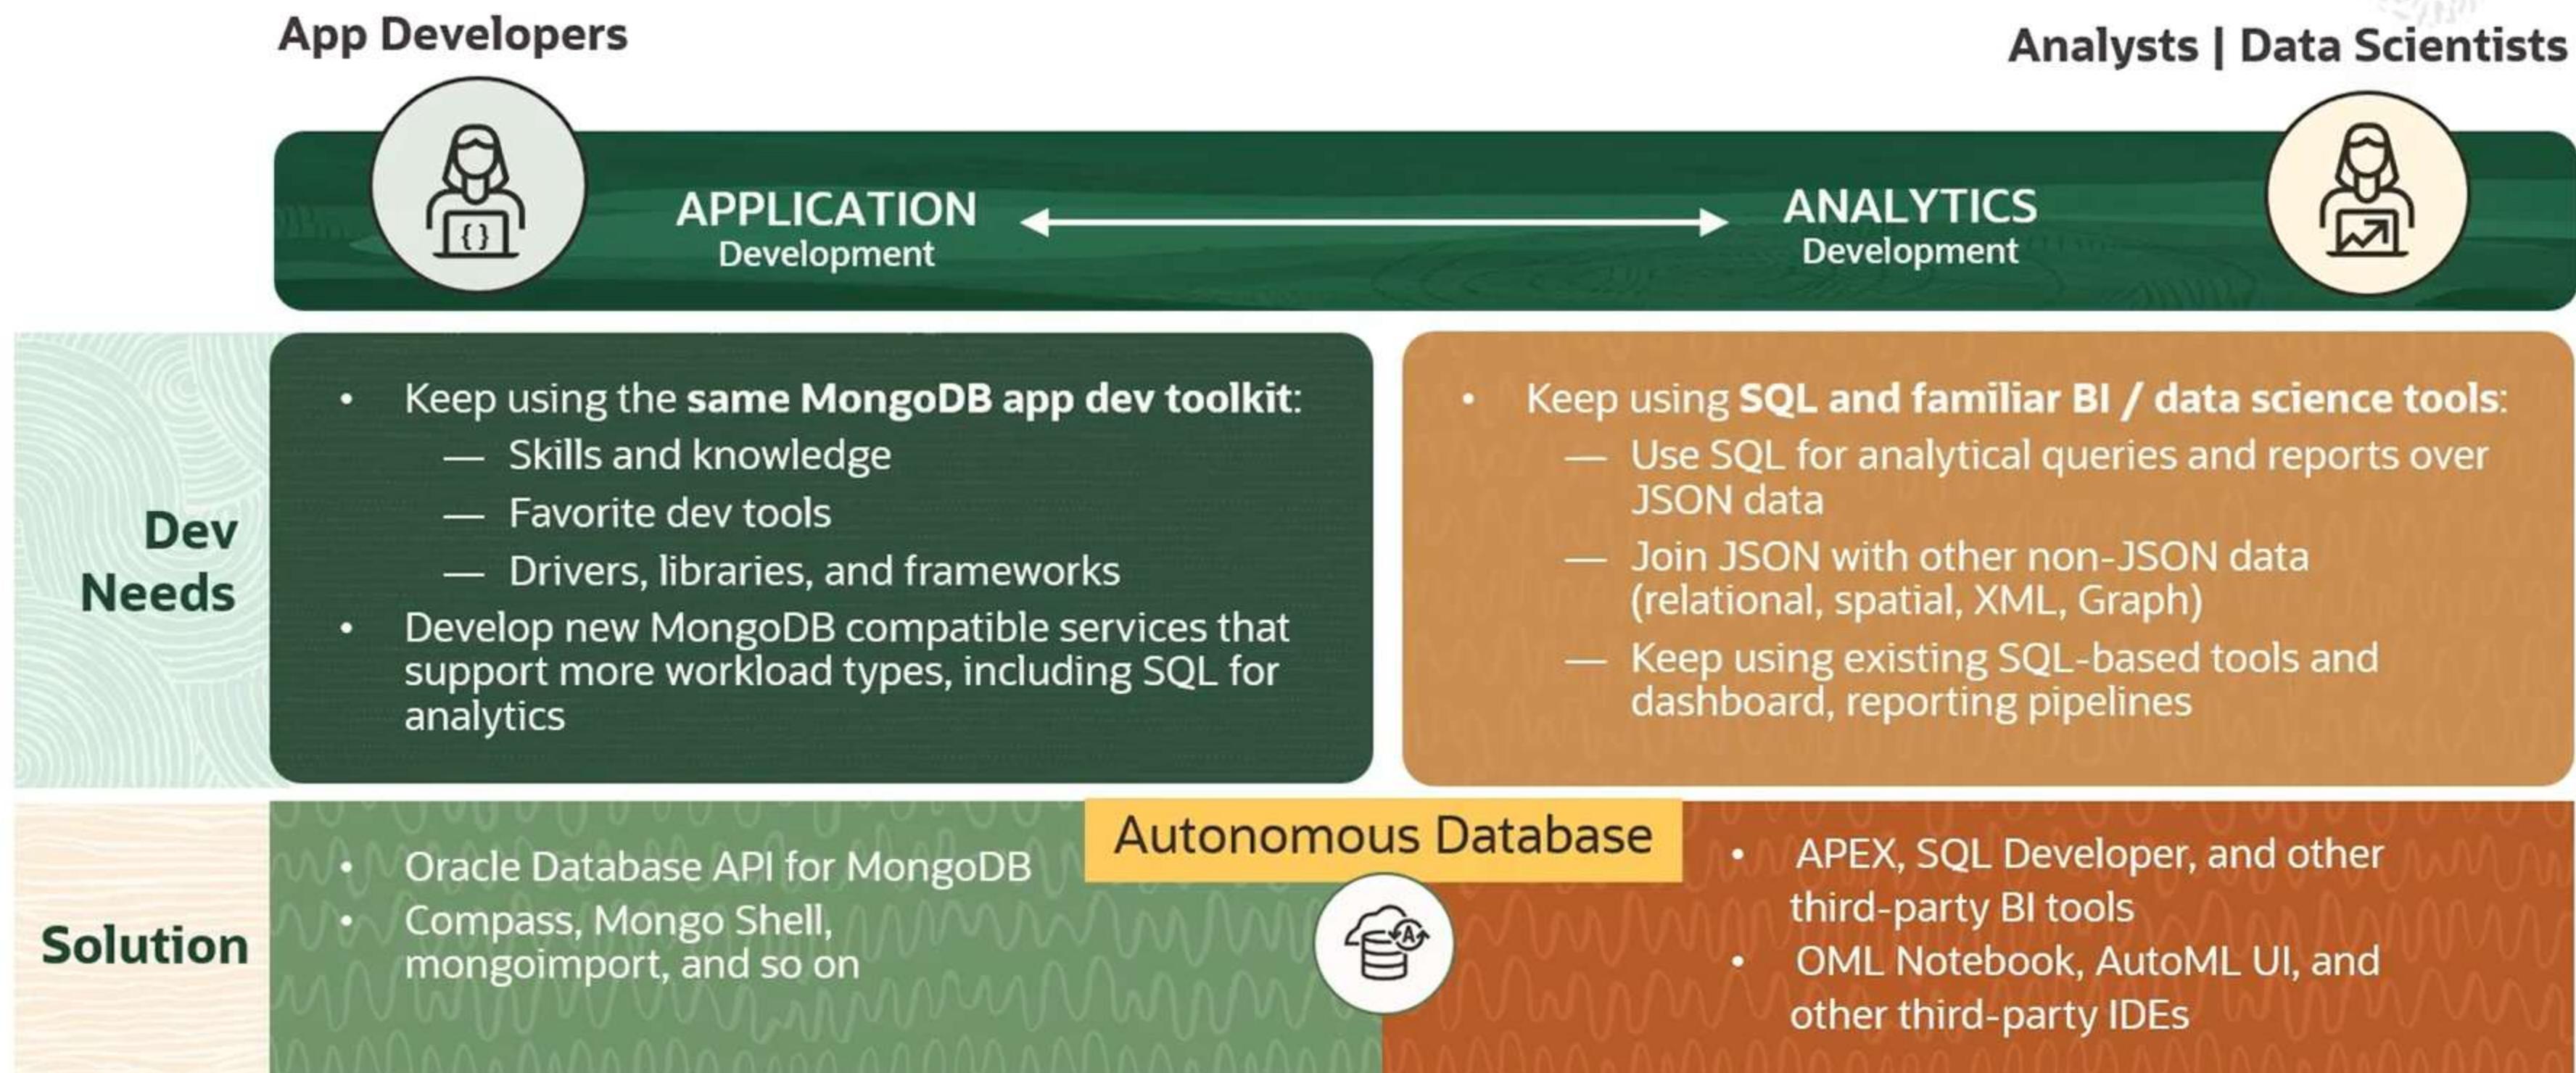
*Simple, flexible persistence
for applications, microservices*

*Powerful analytics and reporting
directly over collections*



Autonomous Database Is MongoDB Compatible* and More

Run MongoDB workload on Autonomous Database with Oracle Database API for MongoDB





Use Oracle Autonomous JSON Database

